

Handling Missing Data

```
In [ ]: important as presence of missing values may lead to error (If None values are present), wrong result (if NAN values are  
2
```

1 ## Missing Data in Pandas

```
2  
3 In Pandas both None and NAN are treated as the indicators of missing data
```

None : Pythonic missing data

The first sentinel value used by Pandas is `None`, a Python singleton object that is often used for missing data in Python code. Because it is a Python object, `None` cannot be used in any arbitrary NumPy/Pandas array, but only in arrays with data type `'object'` (i.e., arrays of Python objects):

```
In [4]: 1 import numpy as np  
2 import pandas as pd
```

```
In [5]: 1 vals1 = np.array([1, None, 3, 4])  
2 vals1
```

```
Out[5]: array([1, None, 3, 4], dtype=object)
```

```
In [7]: 1 vals1.sum() ###Gives an error
        2 ###Bcz you cannot perform any operation with None type
```

```
-----
TypeError                                Traceback (most recent call last)
<ipython-input-7-b9ecc17abfdf> in <module>
----> 1 vals1.sum() ###Gives an error
      2 ###Bcz you can not perform any operation with None type

C:\ProgramData\Anaconda3\lib\site-packages\numpy\core\_methods.py in _sum(a, axis, dtype, out, keepdims, initial)
    34 def _sum(a, axis=None, dtype=None, out=None, keepdims=False,
    35         initial=_NoValue):
----> 36     return umr_sum(a, axis, dtype, out, keepdims, initial)
    37
    38 def _prod(a, axis=None, dtype=None, out=None, keepdims=False,
```

TypeError: unsupported operand type(s) for +: 'int' and 'NoneType'

This reflects the fact that addition between an integer and `None` is undefined.

NaN : Missing numerical data

The other missing data representation, `NaN` (acronym for *Not a Number*), is different; it is a special floating-point value recognized by all systems that use the standard IEEE floating-point representation:

```
In [25]: 1 vals2 = np.array([1, np.nan, 3, 4])
        2 vals2.dtype
```

Out[25]: dtype('float64')

Notice that NumPy chose a native floating-point type for this array: this means that unlike the object array from before, this array supports fast operations pushed into compiled code. You should be aware that `NaN` is a bit like a data virus—it infects any other object it touches. Regardless of the operation, the result of arithmetic with `NaN` will be another `NaN` :

```
In [26]: 1 1 + np.nan
```

```
Out[26]: nan
```

```
In [27]: 1 0 * np.nan
```

```
Out[27]: nan
```

NaN and None in Pandas

NaN and None both have their place, and Pandas is built to handle the two of them nearly interchangeably, converting between them where appropriate:

For types that don't have an available sentinel value, Pandas automatically type-casts when NA values are present. For example, if we set a value in an integer array to `np.nan`, it will automatically be upcast to a floating-point type to accommodate the NA:

```
In [ ]: 1 import pandas as pd
        2 import numpy as np
        3 pd.Series([1, np.nan, 2, None])
```

```
In [8]: 1 x = pd.Series(range(2), dtype=int)
        2 x
```

```
Out[8]: 0    0
        1    1
        dtype: int32
```

```
In [9]: 1 x[0] = None
        2 x
```

```
Out[9]: 0    NaN
        1    1.0
        dtype: float64
```

Notice that in addition to casting the integer array to floating point, Pandas automatically converts the `None` to a `NaN` value. (Be aware that there is a proposal to add a native integer NA to Pandas in the future; as of this writing, it has not been included).

While this type of magic may feel a bit hackish compared to the more unified approach to NA values in domain-specific languages like R, the Pandas sentinel/casting approach works quite well in practice and in my experience only rarely causes issues.

The following table lists the upcasting conventions in Pandas when NA values are introduced:

Typeclass	Conversion When Storing NAs	NA Sentinel Value
floating	No change	np.nan
object	No change	None or np.nan
integer	Cast to float64	np.nan
boolean	Cast to object	None or np.nan

Keep in mind that in Pandas, string data is always stored with an `object` dtype.

Operating on Null Values

As we have seen, Pandas treats `None` and `NaN` as essentially interchangeable for indicating missing or null values. To facilitate this convention, there are several useful methods for detecting, removing, and replacing null values in Pandas data structures. They are:

- `isnull()` : Generate a boolean mask indicating missing values
- `notnull()` : Opposite of `isnull()`
- `dropna()` : Return a filtered version of the data
- `fillna()` : Return a copy of the data with missing values filled or imputed

We will conclude this section with a brief exploration and demonstration of these routines.

Detecting null values

Pandas data structures have two useful methods for detecting null data: `isnull()` and `notnull()`. Either one will return a Boolean mask over the data. For example:

```
In [10]: 1 data = pd.Series([1, np.nan, 'hello', None])
```

```
In [11]: 1 data.isnull()
```

```
Out[11]: 0    False
          1     True
          2    False
          3     True
          dtype: bool
```

As mentioned in [Data Indexing and Selection \(03.02-Data-Indexing-and-Selection.ipynb\)](#), Boolean masks can be used directly as a `Series` or `DataFrame` index:

```
In [12]: 1 data[data.notnull()]
```

```
Out[12]: 0      1
          2   hello
          dtype: object
```

The `isnull()` and `notnull()` methods produce similar Boolean results for `DataFrame` s.

Dropping null values

In addition to the masking used before, there are the convenience methods, `dropna()` (which removes NA values) and `fillna()` (which fills in NA values). For a `Series` , the result is straightforward:

```
In [13]: 1 data.dropna()
```

```
Out[13]: 0      1
          2   hello
          dtype: object
```

For a `DataFrame` , there are more options. Consider the following `DataFrame` :

```
In [14]: 1 df = pd.DataFrame([[1,      np.nan, 2],
2                        [2,      3,    5],
3                        [np.nan, 4,    6]])
4 df
```

Out[14]:

	0	1	2
0	1.0	NaN	2
1	2.0	3.0	5
2	NaN	4.0	6

We cannot drop single values from a `DataFrame` ; we can only drop full rows or full columns. Depending on the application, you might want one or the other, so `dropna()` gives a number of options for a `DataFrame` .

By default, `dropna()` will drop all rows in which *any* null value is present:

```
In [15]: 1 df.dropna()
```

Out[15]:

	0	1	2
1	2.0	3.0	5

Alternatively, you can drop NA values along a different axis; `axis=1` drops all columns containing a null value:

```
In [16]: 1
          2 df.dropna(axis='columns')
```

Out[16]:

```
      2
0 2
1 5
2 6
```

But this drops some good data as well; you might rather be interested in dropping rows or columns with *all* NA values, or a majority of NA values. This can be specified through the `how` or `thresh` parameters, which allow fine control of the number of nulls to allow through.

The default is `how='any'`, such that any row or column (depending on the `axis` keyword) containing a null value will be dropped. You can also specify `how='all'`, which will only drop rows/columns that are *all* null values:

```
In [17]: 1 df[3] = np.nan
          2 df
```

Out[17]:

```
      0    1  2    3
0  1.0  NaN  2  NaN
1  2.0   3.0  5  NaN
2  NaN   4.0  6  NaN
```

```
In [18]: 1 df.dropna(axis='columns', how='all') ###Check the result the third column is dropped (removed)
```

Out[18]:

```
      0    1  2
0  1.0  NaN  2
1  2.0   3.0  5
2  NaN   4.0  6
```

For finer-grained control, the `thresh` parameter lets you specify a minimum number of non-null values for the row/column to be kept:

```
In [19]: 1 df.dropna(axis='rows', thresh=3)
```

Out[19]:

	0	1	2	3
1	2.0	3.0	5	NaN

Here the first and last row have been dropped, because they contain only two non-null values.

Filling null values

Sometimes rather than dropping NA values, you'd rather replace them with a valid value. This value might be a single number like zero, or it might be some sort of imputation or interpolation from the good values. You could do this in-place using the `isnull()` method as a mask, but because it is such a common operation Pandas provides the `fillna()` method, which returns a copy of the array with the null values replaced.

Consider the following Series :

```
In [20]: 1 data = pd.Series([1, np.nan, 2, None, 3], index=list('abcde'))
         2 data
```

Out[20]:

a	1.0
b	NaN
c	2.0
d	NaN
e	3.0

dtype: float64

We can fill NA entries with a single value, such as zero or any number:


```
In [21]: 1 print(data.fillna(0))  
        2 data.fillna(8)
```

```
a    1.0  
b    0.0  
c    2.0  
d    0.0  
e    3.0  
dtype: float64
```

```
Out[21]: a    1.0  
        b    8.0  
        c    2.0  
        d    8.0  
        e    3.0  
        dtype: float64
```

We can specify a forward-fill to propagate the previous value forward:

```
In [22]: 1 # forward-fill  
        2 data.fillna(method='ffill')
```

```
Out[22]: a    1.0  
        b    1.0  
        c    2.0  
        d    2.0  
        e    3.0  
        dtype: float64
```

Or we can specify a back-fill to propagate the next values backward:

```
In [23]: 1 # back-fill  
        2 data.fillna(method='bfill')
```

```
Out[23]: a    1.0  
        b    2.0  
        c    2.0  
        d    3.0  
        e    3.0  
        dtype: float64
```